

Using AI at Work

What to know before you paste work into an AI chat tool.



What we'll cover



01 A few terms first

02 Two examples of AI at work

03 Your account type decides what happens to your data

04 Before you automate: is AI in the loop?

05 Some examples of how to use AI

06 Examples of what not to share

07 Why AI makes things up

08 Who owns what it writes?

A few terms first, so we're all on the same page



1 AI model (or “LLM”)

The engine behind the AI chat tools out there. It's been trained on a huge pile of text to guess a good next answer. It doesn't know things the way a person does.

2 Prompt

Just what you type in: the question or instruction you give it. A better prompt gets you a better answer, which we'll come back to.

3 Training data

Everything a model has learned from. On some accounts, what you type today can end up shaping tomorrow's version of the model, more on that shortly.

4 AI agent

A tool built on top of a model that can actually go do something (search, calculate, pull a file) instead of just chatting back at you.

5 AI hallucination

When it states something false with total confidence: a rule, a source that simply doesn't exist. This is the one to remember: always check anything factual before you rely on it.

Two examples of AI at work



Not the only kinds out there, but a useful line to know.

1 Your company's own AI tool

Something the company sets up and controls itself. It might run on the same kind of model as any public AI chat tool under the hood, but it lives inside the company's own systems, with its own rules about what it can see and where the data goes.

Examples: an assistant your IT team set up, or AI built into tools your company already manages.

2 An outside AI tool

Any AI chat tool you've signed up for yourself through a browser or an app. Genuinely useful. But the moment you're in it, you're a guest inside someone else's product, on their terms.

Examples: a public AI chatbot on your own login, an AI app on your phone, any free browser extension that "adds AI" to something.

Same kind of tool. Completely different rules about what happens to what you type.

Your account type decides what happens to your data



Free, paid personal, and business accounts are three different contracts, not three price points on the same thing.

	Personal (Free)	Personal (Paid)	Business / Enterprise
Trains the model on what you type?	On by default, opt out yourself	On by default, opt out yourself	No, by default
Who can actually see it	The provider	The provider	Your company; provider access is limited by contract
Data deletion on request	Rarely	Sometimes	Yes, configurable
Can the company manage the account?	No	No	Admin controls, company sign-in

A paid personal subscription buys you a faster, better tool. It doesn't buy you privacy.

Disclaimer: data policies like these change constantly and vary by provider, so check the tool's own current data policy and go into its Data Controls yourself. Don't assume the default is doing what you want.

Before you automate: is AI actually in the loop?



Not every automation needs what's on the last page. The ones with an AI step do.

1


No AI involved

A scheduled export, a template macro, a straightforward rules engine moving data from one place to another. Normal build, test, and sign-off process applies, same as any other process change.

2

AI is one of the steps

If any step reads, summarises, decides, or drafts using an AI model, it inherits the same account rules from the last page. Test it on made-up data first, and write down what each piece of that made-up data stands in for.



Some examples of how to use AI



None of this needs code, and none of it needs a new system.

1

Explaining an issue

Describe what's going wrong in plain words and ask what usually causes that pattern.

2

Checking your writing

Ask it to find the gaps or the unclear bit before you send an email or a note.

3

Explaining new things

Market terms, a system, contract language: explained in plain English, no client data needed.



What generally counts as sensitive



1

Client & personal data

Names, employers, contact details: anyone identifiable, colleague or client.

2

Account & case numbers

Account IDs, case references, policy numbers: anything that maps back to a specific record.

3

Exact financial figures

Real balances, amounts, dates or valuations: round them or describe them instead.

4

Credentials & internal info

Logins, passwords, API keys, and anything not yet public: deals, strategy, unreleased results.

!

Rule of thumb

If you wouldn't post it on the company intranet for everyone to see, don't paste it into an AI chat tool either.

Examples of what not to share



NEVER PASTE THIS

Client: Jane Doe, Acme Holdings Ltd
Account: 88-4471-02
Amount: EUR 41,220, due 14 Mar
Login: j.doe /

Names a person, a company, an account number, an amount, and a login. Once it's sent, it can't be pulled back.



Take out the who, keep the what. Names, account numbers and exact amounts come out; the actual problem stays in.

Why AI makes things up (and how to reduce it)



This matters most with an outside tool: there's no one behind it checking your specific answer.

WHY IT HAPPENS

- 1 It can't see your company's systems, only what you tell it.
- 2 It has a cutoff date, a new regulation or policy update might not be in there yet.
- 3 It's built to sound fluent, not to flag doubt: the tone is the same whether it's sure or just making something up.
- 4 Disagreeing with it often makes it agree with you, which isn't the same as you being right.

HOW TO REDUCE IT: PROMPT FOR IT

Try asking it this:

"Always tell me where this comes from. If you don't have a real source, or you're not sure, say so instead of guessing."

Who owns what it writes?



Ownership of AI-written material is still an open legal question, treat it that way.

THE SHORT VERSION

Courts have been ruling that something purely AI-written isn't protected by copyright at all, only the parts where a person actually made the creative decisions. Typing a good prompt, or picking your favorite draft out of a few, doesn't count as authorship on its own.

In practice: a report or client note that's mostly untouched AI output may not be something the company can fully claim as its own work, or stop someone else from copying.

HOW TO STAY COVERED

1

Edit it, don't just copy it.

Real changes and choices are what make something actually yours.

2

Keep a record.

Note what was AI-drafted and what a person changed, especially for anything client-facing.

3

Don't publish it as-is.

Not under the company's name, for anything that matters legally or commercially.

4

Ask when it's not obvious.

Check your company's AI policy, or ask your legal team, before your name goes on AI-written material.

A black and white photograph showing a robotic hand on the left and a human hand on the right. The two hands are positioned so that their fingers meet at the top and bottom, forming a heart shape. The robotic hand is dark and metallic, with visible joints and segments. The human hand is light-skinned and appears to be a right hand. The background is plain white.

*If you were wondering, yes, this
presentation was made with AI
help.*

Explore the data behind this guide
sansanabria.github.io/ai-adoption-eu